

1 Environment

I trained this model on a Windows 11 machine with an NVIDIA RTX 5080 GPU using Conda with Python 3.9.25 and CUDA 12.8. This open issue has been particularly helpful in resolving Python/CUDA dependency issues: <https://github.com/graphdeco-inria/gaussian-splatting/issues/1215>.

2 Data Split

I utilized the "--eval" flag to automatically split the dataset for evaluation.

3 Metrics

I ran several iterations and the best of which achieved the following metrics.

Metric	Value	Reference Target
PSNR (dB)	22.797117	≥ 20.0
SSIM	0.8152919	≥ 0.80
LPIPS	0.1899525	$\leq ?$

4 Screenshots

5 Analysis

Training took, on average, around 6 minutes and 30 seconds to complete at 30,000 iterations. Using the "--optimizer_type sparse_adam" flag brought the average training time down to around 3 minutes.