

ECEN 743: Reinforcement Learning

Policy Gradient Theorem

Dileep Kalathil
Texas A&M University

References

- [SB, Chapter 13]
- [MMT, Chapter 12]
- [AJKS, Section 3]

Reinforcement Learning Questions

- How do we learn the value of a policy π ?
- How do we learn the optimal action-value function Q^* ?
- How do we learn the optimal policy π^* ?

... without the knowledge of the model P

- Need to **learn** from the observed sequence of states, actions, and rewards

Finding the Optimal Policy Parameter

- Value of a policy; $V_{\pi}(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$

Finding the Optimal Policy Parameter

- Value of a policy; $V_{\pi}(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0

Finding the Optimal Policy Parameter

- Value of a policy; $V_{\pi}(s) = \mathbb{E} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_{\pi}(s)]$

Finding the Optimal Policy Parameter

- Value of a policy; $V_\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_\pi(s)]$
 - ▶ We will often just use V_π , making the dependence on μ_0 implicit

Finding the Optimal Policy Parameter

- Value of a policy; $V_\pi(s) = \mathbb{E} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_\pi(s)]$
 - ▶ We will often just use V_π , making the dependence on μ_0 implicit
- Consider the parameterized policies π_θ , where $\theta \in \Theta \subset \mathbb{R}^d$

Finding the Optimal Policy Parameter

- Value of a policy; $V_\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_\pi(s)]$
 - ▶ We will often just use V_π , making the dependence on μ_0 implicit
- Consider the parameterized policies π_θ , where $\theta \in \Theta \subset \mathbb{R}^d$
- V_{π_θ} is the value of the policy π_θ

Finding the Optimal Policy Parameter

- Value of a policy; $V_{\pi}(s) = \mathbb{E} [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_{\pi}(s)]$
 - ▶ We will often just use V_{π} , making the dependence on μ_0 implicit
- Consider the parameterized policies π_{θ} , where $\theta \in \Theta \subset \mathbb{R}^d$
- $V_{\pi_{\theta}}$ is the value of the policy π_{θ}
- Policy optimization problem:

$$\max_{\theta \in \Theta} V_{\pi_{\theta}}$$

Finding the Optimal Policy Parameter

- Value of a policy; $V_\pi(s) = \mathbb{E}[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) | s_0 = s]$
- Assume that the initial state s_0 is sampled according to a distribution μ_0
- Define $V_{\pi, \mu_0} = \mathbb{E}_{s \sim \mu_0(\cdot)} [V_\pi(s)]$
 - ▶ We will often just use V_π , making the dependence on μ_0 implicit
- Consider the parameterized policies π_θ , where $\theta \in \Theta \subset \mathbb{R}^d$
- V_{π_θ} is the value of the policy π_θ
- Policy optimization problem:

$$\max_{\theta \in \Theta} V_{\pi_\theta}$$

- Can we develop (stochastic) gradient algorithms for solving the above problem?

Policy Parameterization

- Tabular representation

$$\pi(s, a) = \theta_{s,a}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

Policy Parameterization

- Tabular representation

$$\pi(s, a) = \theta_{s,a}, \quad \theta \in \mathbb{R}^{|S||\mathcal{A}|}$$

- Softmax policies for tabular representation

$$\pi(s, a) = \frac{\exp(\theta_{s,a})}{\sum_b \exp(\theta_{s,b})}, \quad \theta \in \mathbb{R}^{|S||\mathcal{A}|}$$

Policy Parameterization

- Tabular representation

$$\pi(s, a) = \theta_{s,a}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies for tabular representation

$$\pi(s, a) = \frac{\exp(\theta_{s,a})}{\sum_b \exp(\theta_{s,b})}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies / log-linear policies

Consider the basis function $\phi(s, a) \in \mathbb{R}^d$, $d \ll |\mathcal{S}||\mathcal{A}|$. Let $\theta \in \mathbb{R}^d$. Then,

$$\pi(s, a) = \frac{\exp(\theta^\top \phi(s, a))}{\sum_b \exp(\theta^\top \phi(s, b))}$$

Policy Parameterization

- Tabular representation

$$\pi(s, a) = \theta_{s,a}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies for tabular representation

$$\pi(s, a) = \frac{\exp(\theta_{s,a})}{\sum_b \exp(\theta_{s,b})}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies / log-linear policies

Consider the basis function $\phi(s, a) \in \mathbb{R}^d$, $d \ll |\mathcal{S}||\mathcal{A}|$. Let $\theta \in \mathbb{R}^d$. Then,

$$\pi(s, a) = \frac{\exp(\theta^\top \phi(s, a))}{\sum_b \exp(\theta^\top \phi(s, b))}$$

- Neural softmax policies

$$\pi(s, a) = \frac{\exp(f_\theta(s, a))}{\sum_b \exp(f_\theta(s, b))}$$

Policy Parameterization

- Tabular representation

$$\pi(s, a) = \theta_{s,a}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies for tabular representation

$$\pi(s, a) = \frac{\exp(\theta_{s,a})}{\sum_b \exp(\theta_{s,b})}, \quad \theta \in \mathbb{R}^{|\mathcal{S}||\mathcal{A}|}$$

- Softmax policies / log-linear policies

Consider the basis function $\phi(s, a) \in \mathbb{R}^d$, $d \ll |\mathcal{S}||\mathcal{A}|$. Let $\theta \in \mathbb{R}^d$. Then,

$$\pi(s, a) = \frac{\exp(\theta^\top \phi(s, a))}{\sum_b \exp(\theta^\top \phi(s, b))}$$

- Neural softmax policies

$$\pi(s, a) = \frac{\exp(f_\theta(s, a))}{\sum_b \exp(f_\theta(s, b))}$$

- Gaussian policies

$$\pi(s, a) = \mathcal{N}(\mu_\theta(s), \Sigma_\theta(s))$$

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$
- But, the gradient of the (stationary) distribution of the states induced by the policy cannot be computed easily

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$
- But, the gradient of the (stationary) distribution of the states induced by the policy cannot be computed easily
 - ▶ Requires the knowledge of the transition probability, which is unknown

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$
- But, the gradient of the (stationary) distribution of the states induced by the policy cannot be computed easily
 - ▶ Requires the knowledge of the transition probability, which is unknown
- How do we estimate the gradient?

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$
- But, the gradient of the (stationary) distribution of the states induced by the policy cannot be computed easily
 - ▶ Requires the knowledge of the transition probability, which is unknown
- How do we estimate the gradient?

Gradient of the Value Function

- Stochastic gradient ascent for learning the optimal policy parameter

$$\theta \leftarrow \theta + \alpha \widehat{\nabla_{\theta} V_{\pi_{\theta}}}$$

- Computing $\nabla_{\theta} V_{\pi_{\theta}}$ is challenging
- Value of a policy depends on both action selections (policy) and the (stationary) distribution of the states resulting from the policy
- We can directly compute $\nabla_{\theta} \pi_{\theta}$
- But, the gradient of the (stationary) distribution of the states induced by the policy cannot be computed easily
 - ▶ Requires the knowledge of the transition probability, which is unknown
- How do we estimate the gradient?

Policy Gradient Theorem!

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$P_{\text{traj}, \theta}(\tau) = \mathbb{P}(s_T, a_T, \tau[0 : T - 1]) = \mathbb{P}(s_T, a_T | \tau[0 : T - 1]) P_{\text{traj}, \theta}(\tau[0 : T - 1])$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$\begin{aligned} P_{\text{traj}, \theta}(\tau) &= \mathbb{P}(s_T, a_T, \tau[0 : T-1]) = \mathbb{P}(s_T, a_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \mathbb{P}(a_T | s_T, \tau[0 : T-1]) \mathbb{P}(s_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \end{aligned}$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$\begin{aligned} P_{\text{traj}, \theta}(\tau) &= \mathbb{P}(s_T, a_T, \tau[0 : T-1]) = \mathbb{P}(s_T, a_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \mathbb{P}(a_T | s_T, \tau[0 : T-1]) \mathbb{P}(s_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(\tau[0 : T-1]) \end{aligned}$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$\begin{aligned} P_{\text{traj}, \theta}(\tau) &= \mathbb{P}(s_T, a_T, \tau[0 : T-1]) = \mathbb{P}(s_T, a_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \mathbb{P}(a_T | s_T, \tau[0 : T-1]) \mathbb{P}(s_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(s_{T-1}, a_{T-1}, \tau[0 : T-2]) \end{aligned}$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$\begin{aligned} P_{\text{traj}, \theta}(\tau) &= \mathbb{P}(s_T, a_T, \tau[0 : T-1]) = \mathbb{P}(s_T, a_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \mathbb{P}(a_T | s_T, \tau[0 : T-1]) \mathbb{P}(s_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(s_{T-1}, a_{T-1}, \tau[0 : T-2]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) \pi_\theta(s_{T-1}, a_{T-1}) P(s_{T-1} | s_{T-2}, a_{T-2}) P_{\text{traj}, \theta}(\tau[0 : T-2]) \end{aligned}$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

$$\begin{aligned} P_{\text{traj}, \theta}(\tau) &= \mathbb{P}(s_T, a_T, \tau[0 : T-1]) = \mathbb{P}(s_T, a_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \mathbb{P}(a_T | s_T, \tau[0 : T-1]) \mathbb{P}(s_T | \tau[0 : T-1]) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(\tau[0 : T-1]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) P_{\text{traj}, \theta}(s_{T-1}, a_{T-1}, \tau[0 : T-2]) \\ &= \pi_\theta(s_T, a_T) P(s_T | s_{T-1}, a_{T-1}) \pi_\theta(s_{T-1}, a_{T-1}) P(s_{T-1} | s_{T-2}, a_{T-2}) P_{\text{traj}, \theta}(\tau[0 : T-2]) \\ &= \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t) \end{aligned}$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

Probability of a Trajectory

- Let τ be a trajectory generated using π_θ ,

$$\tau = (s_0, a_0, s_1, a_1, \dots, s_T, a_T), \quad a_t \sim \pi_\theta(s_t, \cdot), s_{t+1} \sim P(\cdot | s_t, a_t), s_0 \sim \mu_0(\cdot)$$

- What is the probability of the trajectory τ ?

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t=0}^T \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

- We can extend this to arbitrary long trajectories, $\tau = (s_t, a_t)_{t \geq 0}$,

$$P_{\text{traj}, \theta}(\tau) = \mu_0(s_0) \prod_{t \geq 0} \pi_\theta(s_t, a_t) P(s_{t+1} | s_t, a_t)$$

Gradient of the Value Function: Form 1

- Let $\tau = (s_t, a_t)_{t \geq 0}$, $s_0 \sim \mu_0(\cdot)$. Denote the cumulative discounted reward of trajectory τ as

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

Gradient of the Value Function: Form 1

- Let $\tau = (s_t, a_t)_{t \geq 0}$, $s_0 \sim \mu_0(\cdot)$. Denote the cumulative discounted reward of trajectory τ as

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

- It is straight forward to note that

$$V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)]$$

Gradient of the Value Function: Form 1

- Let $\tau = (s_t, a_t)_{t \geq 0}$, $s_0 \sim \mu_0(\cdot)$. Denote the cumulative discounted reward of trajectory τ as

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

- It is straight forward to note that

$$V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)]$$

Theorem (PG: Form 1)

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

Gradient of the Value Function: Form 1

- Let $\tau = (s_t, a_t)_{t \geq 0}$, $s_0 \sim \mu_0(\cdot)$. Denote the cumulative discounted reward of trajectory τ as

$$R(\tau) = \sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)$$

- It is straight forward to note that

$$V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)]$$

Theorem (PG: Form 1)

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- Gradient above does not depend on the gradient of the (stationary) distribution induced by the policy!

Proof of Policy Gradient Theorem: Form 1

Proof:



Proof of Policy Gradient Theorem: Form 1

Proof:



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\nabla V_{\pi_{\theta}} = \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau)$$



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \\ &= \sum_{\tau} R(\tau) \nabla P_{\text{traj}, \theta}(\tau) = \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log P_{\text{traj}, \theta}(\tau)\end{aligned}$$



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \\ &= \sum_{\tau} R(\tau) \nabla P_{\text{traj}, \theta}(\tau) = \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log P_{\text{traj}, \theta}(\tau) \\ &= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log \left(\mu_0(s_0) \prod_{t=0}^{\infty} \pi_{\theta}(s_t, a_t) P(s_{t+1} | s_t, a_t) \right)\end{aligned}$$



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) \nabla P_{\text{traj}, \theta}(\tau) = \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log \left(\mu_0(s_0) \prod_{t=0}^{\infty} \pi_{\theta}(s_t, a_t) P(s_{t+1} | s_t, a_t) \right) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \left(\sum_{t=0}^{\infty} \log \pi_{\theta}(s_t, a_t) \right)\end{aligned}$$



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) \nabla P_{\text{traj}, \theta}(\tau) = \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log \left(\mu_0(s_0) \prod_{t=0}^{\infty} \pi_{\theta}(s_t, a_t) P(s_{t+1} | s_t, a_t) \right) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \left(\sum_{t=0}^{\infty} \log \pi_{\theta}(s_t, a_t) \right) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \left(\sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t) \right)\end{aligned}$$



Proof of Policy Gradient Theorem: Form 1

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) \nabla P_{\text{traj}, \theta}(\tau) = \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log P_{\text{traj}, \theta}(\tau) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \log \left(\mu_0(s_0) \prod_{t=0}^{\infty} \pi_{\theta}(s_t, a_t) P(s_{t+1} | s_t, a_t) \right) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \nabla \left(\sum_{t=0}^{\infty} \log \pi_{\theta}(s_t, a_t) \right) \\&= \sum_{\tau} R(\tau) P_{\text{traj}, \theta}(\tau) \left(\sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t) \right) \\&= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t) \right]\end{aligned}$$



REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$
 - ▶ For each trajectory τ_i , get $[R(\tau_i) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t^i, a_t^i)]$

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$
 - ▶ For each trajectory τ_i , get $[R(\tau_i) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t^i, a_t^i)]$
 - ▶ Average to get $\widehat{\nabla V_{\pi_{\theta}}}$

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$
 - ▶ For each trajectory τ_i , get $[R(\tau_i) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t^i, a_t^i)]$
 - ▶ Average to get $\widehat{\nabla V_{\pi_{\theta}}}$
- This approach, however, is unreliable due very high variance

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$
 - ▶ For each trajectory τ_i , get $[R(\tau_i) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t^i, a_t^i)]$
 - ▶ Average to get $\widehat{\nabla V_{\pi_{\theta}}}$
- This approach, however, is unreliable due very high variance
- This issue is avoided by exploiting the temporal structure and introducing a baseline

REINFORCE: Monte Carlo Policy Gradient

- Consider a finite horizon (episodic) RL problem
- We have $R(\tau) = \sum_{t=0}^T \gamma^t r(s_t, a_t)$, and

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

- How do we get an estimate $\widehat{\nabla V_{\pi_{\theta}}}$?
 - ▶ Generate n trajectories $(\tau_i)_{i=1}^n$
 - ▶ For each trajectory τ_i , get $[R(\tau_i) \sum_{t=0}^T \nabla \log \pi_{\theta}(s_t^i, a_t^i)]$
 - ▶ Average to get $\widehat{\nabla V_{\pi_{\theta}}}$
- This approach, however, is unreliable due very high variance
- This issue is avoided by exploiting the temporal structure and introducing a baseline
 - ▶ Variance reduction for PG is an active area of research

Using Temporal Structure

- For ease of notations, let us assume a finite T and $\gamma = 1$

Using Temporal Structure

- For ease of notations, let us assume a finite T and $\gamma = 1$
- We have:

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[R(\tau)] = \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}\left[\sum_{t=1}^T r_t\right] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[R(\tau)] \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t)]$$

Using Temporal Structure

- For ease of notations, let us assume a finite T and $\gamma = 1$
- We have:

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[R(\tau)] = \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}\left[\sum_{t=1}^T r_t\right] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[R(\tau)] \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t)]$$

- Using the same derivation as before, we can show that

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[r_{t'}] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)}[r_{t'}] \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$$

Using Temporal Structure

- For ease of notations, let us assume a finite T and $\gamma = 1$
- We have:

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=1}^T r_t \right] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] \sum_{t=0}^{\infty} \nabla \log \pi_{\theta}(s_t, a_t)]$$

- Using the same derivation as before, we can show that

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [r_{t'}] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$$

- Summing over t' , we get

$$\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

Using Temporal Structure

- We have: $\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$

Using Temporal Structure

- We have: $\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$
- By exchanging the order of the summation, we can rewrite the RHS as:

$$\mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t) \right] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \sum_{t'=t}^T r_{t'} \right]$$

Using Temporal Structure

- We have: $\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$
- By exchanging the order of the summation, we can rewrite the RHS as:

$$\mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \sum_{t'=t}^T r_{t'}]$$

- Recall the definition of **Return**: $G_t = \sum_{t'=t}^T r_{t'}$.

Using Temporal Structure

- We have: $\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$
- By exchanging the order of the summation, we can rewrite the RHS as:

$$\mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \sum_{t'=t}^T r_{t'}]$$

- Recall the definition of **Return**: $G_t = \sum_{t'=t}^T r_{t'}$.
- So, by leveraging this temporal structural, we get:

Using Temporal Structure

- We have: $\nabla \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [R(\tau)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)]$
- By exchanging the order of the summation, we can rewrite the RHS as:

$$\mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t'=1}^T r_{t'} \sum_{t=0}^{t'} \nabla \log \pi_{\theta}(s_t, a_t)] = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} [\sum_{t=0}^T \nabla \log \pi_{\theta}(s_t, a_t) \sum_{t'=t}^T r_{t'}]$$

- Recall the definition of **Return**: $G_t = \sum_{t'=t}^T r_{t'}$.
- So, by leveraging this temporal structural, we get:

Theorem (PG: Form 2)

$$\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_{\theta}(s_t, a_t) \right]$$

REINFORCE using Temporal Structure

MC Estimate to Trained Estimate

- We have: $\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_{\theta}(s_t, a_t) \right]$

MC Estimate to Trained Estimate

- We have: $\nabla V_{\pi_{\theta}} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_{\theta}(s_t, a_t) \right]$

$$\begin{aligned} \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_{\theta}(s_t, a_t) \right] &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T \mathbb{E}[G_t \nabla \log \pi_{\theta}(s_t, a_t) | (s_t, a_t)] \right] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T \mathbb{E}[G_t | (s_t, a_t)] \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \end{aligned}$$

MC Estimate to Trained Estimate

- We have: $\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_\theta(s_t, a_t) \right]$

$$\begin{aligned} \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T G_t \nabla \log \pi_\theta(s_t, a_t) \right] &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T \mathbb{E}[G_t \nabla \log \pi_\theta(s_t, a_t) | (s_t, a_t)] \right] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T \mathbb{E}[G_t | (s_t, a_t)] \nabla \log \pi_\theta(s_t, a_t) \right] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T Q_{\pi_\theta}(s_t, a_t) \nabla \log \pi_\theta(s_t, a_t) \right] \end{aligned}$$

Theorem (PG: Form 3)

$$\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[\sum_{t=0}^T Q_{\pi_\theta}(s_t, a_t) \nabla \log \pi_\theta(s_t, a_t) \right]$$

Visitation Distribution

- For a given policy π , (discounted) state-action visitation distribution is defined as

$$\rho_{\pi, \mu_0}(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a),$$

where $(s_t, a_t)_{t \geq 0}$ is the trajectory generated according to π , and s_0 is sampled according to a given distribution μ_0 .

Visitation Distribution

- For a given policy π , (discounted) state-action visitation distribution is defined as

$$\rho_{\pi, \mu_0}(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a),$$

where $(s_t, a_t)_{t \geq 0}$ is the trajectory generated according to π , and s_0 is sampled according to a given distribution μ_0 .

- We will often just use ρ_{π} , making the dependence on μ_0 implicit.

Visitation Distribution

- For a given policy π , (discounted) state-action visitation distribution is defined as

$$\rho_{\pi, \mu_0}(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a),$$

where $(s_t, a_t)_{t \geq 0}$ is the trajectory generated according to π , and s_0 is sampled according to a given distribution μ_0 .

- We will often just use ρ_{π} , making the dependence on μ_0 implicit.
- Similarly, for a given policy π , (discounted) state visitation distribution is defined as

$$\rho_{\pi}(s) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s),$$

where $(s_t)_{t \geq 0}$ is the state-trajectory generated according to π (we are *overloading* the notation)

Visitation Distribution

- For a given policy π , (discounted) state-action visitation distribution is defined as

$$\rho_{\pi, \mu_0}(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a),$$

where $(s_t, a_t)_{t \geq 0}$ is the trajectory generated according to π , and s_0 is sampled according to a given distribution μ_0 .

- We will often just use ρ_π , making the dependence on μ_0 implicit.
- Similarly, for a given policy π , (discounted) state visitation distribution is defined as

$$\rho_\pi(s) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s),$$

where $(s_t)_{t \geq 0}$ is the state-trajectory generated according to π (we are *overloading* the notation)

- It is straight forward to show that

$$\rho_\pi(x, a) = \rho_\pi(x) \pi(x, a)$$

Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Proof:



Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Proof:

$$V_{\pi} = \mathbb{E}_{\tau \sim P_{\text{traj}, \pi}(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right]$$



Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Proof:

$$\begin{aligned} V_{\pi} &= \mathbb{E}_{\tau \sim P_{\text{traj}, \pi}(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right] \\ &= \sum_s \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t r(s, a) \end{aligned}$$



Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Proof:

$$\begin{aligned} V_{\pi} &= \mathbb{E}_{\tau \sim P_{\text{traj}, \pi}(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right] \\ &= \sum_s \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t r(s, a) \\ &= \sum_s \sum_a r(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \end{aligned}$$



Visitation Distribution and Value Function

Lemma

$$V_{\pi} = \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi}(\cdot, \cdot)} [r(s, a)] = \frac{1}{(1-\gamma)} \sum_{(s,a)} \rho_{\pi}(s, a) r(s, a)$$

Proof:

$$\begin{aligned} V_{\pi} &= \mathbb{E}_{\tau \sim P_{\text{traj}, \pi}(\cdot)} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \right] \\ &= \sum_s \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t r(s, a) \\ &= \sum_s \sum_a r(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \\ &= \frac{1}{(1-\gamma)} \sum_x \sum_a r(s, a) \rho_{\pi}(s, a) \end{aligned}$$



Gradient of the Value Function: Form 3

Theorem (Policy Gradient Theorem - Form 3)

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \frac{1}{(1 - \gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right]\end{aligned}$$

Gradient of the Value Function: Form 3

Theorem (Policy Gradient Theorem - Form 3)

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right]\end{aligned}$$

- How do we get $Q_{\pi_{\theta}}$?

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_{\theta}} = \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)]$

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_{\theta}} = \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)]$
- How do we get $Q_{\pi_{\theta}}(s, a)$?

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_{\theta}} = \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)]$
- How do we get $Q_{\pi_{\theta}}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_{\theta}}(s, a)$
- Then, $\nabla V_{\pi_{\theta}} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_{\theta}(s, a)]$

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$
- Then, $\nabla V_{\pi_\theta} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_\theta(s, a)]$
- Can we then do the update:

$$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla \log \pi_\theta(s, a)$$

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$
- Then, $\nabla V_{\pi_\theta} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_\theta(s, a)]$
- Can we then do the update:

$$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla \log \pi_\theta(s, a)$$

- What is the w to be used?

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$
- Then, $\nabla V_{\pi_\theta} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_\theta(s, a)]$
- Can we then do the update:

$$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla \log \pi_\theta(s, a)$$

- What is the w to be used?
 - ▶ w should give a good approximation of Q -value function for the current policy π_θ

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$
- Then, $\nabla V_{\pi_\theta} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_\theta(s, a)]$
- Can we then do the update:

$$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla \log \pi_\theta(s, a)$$

- What is the w to be used?
 - ▶ w should give a good approximation of Q -value function for the current policy π_θ
 - ▶ Q_{π_θ} is unknown. So, we need to learn w

Actor-Critic Algorithm

- We have, $\nabla V_{\pi_\theta} = \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we get $Q_{\pi_\theta}(s, a)$?
- We can approximate Q -function: $Q_w(s, a) \approx Q_{\pi_\theta}(s, a)$
- Then, $\nabla V_{\pi_\theta} \approx \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_w(s, a) \nabla \log \pi_\theta(s, a)]$
- Can we then do the update:

$$\theta \leftarrow \theta + \alpha Q_w(s, a) \nabla \log \pi_\theta(s, a)$$

- What is the w to be used?
 - ▶ w should give a good approximation of Q -value function for the current policy π_θ
 - ▶ Q_{π_θ} is unknown. So, we need to learn w
 - ▶ θ is changing. So, w should also change with respect to θ

Actor-Critic Algorithm

- **Critic:** Updates Q-value parameter w
 - ▶ (*Evaluate* the policy corresponding to θ)
- **Actor:** Updates policy parameter θ
 - ▶ (*Improve* the policy corresponding to based on the feedback from critic)

Actor-Critic Algorithm

- **Critic:** Updates Q-value parameter w
 - ▶ (*Evaluate* the policy corresponding to θ)
- **Actor:** Updates policy parameter θ
 - ▶ (*Improve* the policy corresponding to based on the feedback from critic)

for each step **do**

$$\theta = \theta + \alpha_{\theta} \nabla_{\theta} \log(\pi_{\theta}(s, a))$$

Sample the next state s' , sample the next action $a' \sim \pi_{\theta}(s', \cdot)$

$$\delta = r(s, a) + \gamma Q_w(s', a') - Q_w(s, a)$$

$$w = w + \alpha_w \delta \nabla_w Q_w(s, a)$$

end for

Actor-Critic Algorithm

- **Critic:** Updates Q-value parameter w
 - ▶ (*Evaluate* the policy corresponding to θ)
- **Actor:** Updates policy parameter θ
 - ▶ (*Improve* the policy corresponding to based on the feedback from critic)

for each step **do**

$$\theta = \theta + \alpha_{\theta} Q_w(s, a) \nabla_{\theta} \log(\pi_{\theta}(s, a))$$

Sample the next state s' , sample the next action $a' \sim \pi_{\theta}(s', \cdot)$

$$\delta = r(s, a) + \gamma Q_w(s', a') - Q_w(s, a)$$

$$w = w + \alpha_w \delta \nabla_w Q_w(s, a)$$

end for

- Learning rate α_{θ} and α_w should be selected appropriately for convergence

Actor-Critic Algorithm

- **Critic:** Updates Q-value parameter w
 - ▶ (*Evaluate* the policy corresponding to θ)
- **Actor:** Updates policy parameter θ
 - ▶ (*Improve* the policy corresponding to based on the feedback from critic)

for each step **do**

$$\theta = \theta + \alpha_{\theta} \nabla_{\theta} \log(\pi_{\theta}(s, a))$$

Sample the next state s' , sample the next action $a' \sim \pi_{\theta}(s', \cdot)$

$$\delta = r(s, a) + \gamma Q_w(s', a') - Q_w(s, a)$$

$$w = w + \alpha_w \delta \nabla_w Q_w(s, a)$$

end for

- Learning rate α_{θ} and α_w should be selected appropriately for convergence
- This is an example of **two timescale stochastic approximation algorithm**

Proof of Policy Gradient Theorem: Form 3

Proof:

Proof of Policy Gradient Theorem: Form 3

Proof:

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\nabla V_{\pi_{\theta}} = \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s)$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\ &= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\ &= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\ &= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\&= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\&= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla Q_{\pi_{\theta}}(s_0, a_0)]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\&= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \\&\quad \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla (r(s_0, a_0) + \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [V_{\pi_{\theta}}(s_1)])]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\&= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \\&\quad \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla (r(s_0, a_0) + \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [V_{\pi_{\theta}}(s_1)])] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \\&\quad \gamma \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [\nabla V_{\pi_{\theta}}(s_1)]\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned}\nabla V_{\pi_{\theta}} &= \sum_s \mathbb{P}(s_0 = s) \nabla V_{\pi_{\theta}}(s) \\&= \nabla \mathbb{E}_{s_0 \sim \mu_0} [V_{\pi_{\theta}}(s_0)] = \mathbb{E}_{s_0 \sim \mu_0} [\nabla \sum_{a_0} \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \nabla \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0) + \sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} [\sum_{a_0} \pi_{\theta}(s_0, a_0) \nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla Q_{\pi_{\theta}}(s_0, a_0)] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \\&\quad \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla (r(s_0, a_0) + \gamma \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [V_{\pi_{\theta}}(s_1)])] \\&= \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} [\nabla \log \pi_{\theta}(s_0, a_0) Q_{\pi_{\theta}}(s_0, a_0)] + \\&\quad \gamma \mathbb{E}_{s_0 \sim \mu_0} \mathbb{E}_{a_0 \sim \pi_{\theta}(s_0, \cdot)} \mathbb{E}_{s_1 \sim P(\cdot | s_0, a_0)} [\nabla V_{\pi_{\theta}}(s_1)] \\&= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_{\theta}(s, a) Q_{\pi_{\theta}}(s, a)] + \gamma \sum_s \mathbb{P}(s_1 = s) \nabla V_{\pi_{\theta}}(s)\end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

Proof of Policy Gradient Theorem: Form 3

Proof:

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \end{aligned}$$



Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \\ &\vdots \end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \\ &\quad \vdots \\ &= \sum_{t=0}^{\infty} \gamma^t \sum_{(s,a)} \mathbb{P}(s_t = s, a_t = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \\ &\quad \vdots \\ &= \sum_{t=0}^{\infty} \gamma^t \sum_{(s,a)} \mathbb{P}(s_t = s, a_t = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &= \sum_{(s,a)} [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(x, a)] \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a) \end{aligned}$$

Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \\ &\quad \vdots \\ &= \sum_{t=0}^{\infty} \gamma^t \sum_{(s,a)} \mathbb{P}(s_t = s, a_t = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &= \sum_{(s,a)} [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a) \\ &= \frac{1}{(1-\gamma)} \sum_{(s,a)} [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \rho_{\pi_\theta}(s, a) \end{aligned}$$



Proof of Policy Gradient Theorem: Form 3

Proof:

$$\begin{aligned} &= \sum_{(s,a)} \mathbb{P}(s_0 = s, a_0 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &+ \gamma \sum_{(s,a)} \mathbb{P}(s_1 = s, a_1 = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] + \gamma^2 \sum_s \mathbb{P}(s_2 = s) \nabla V_{\pi_\theta}(s) \\ &\quad \vdots \\ &= \sum_{t=0}^{\infty} \gamma^t \sum_{(s,a)} \mathbb{P}(s_t = s, a_t = a) [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \\ &= \sum_{(s,a)} [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \sum_{t=0}^{\infty} \gamma^t \mathbb{P}(s_t = s, a_t = a) \\ &= \frac{1}{(1-\gamma)} \sum_{(s,a)} [\nabla \log \pi_\theta(s, a) Q_{\pi_\theta}(s, a)] \rho_{\pi_\theta}(s, a) \\ &= \frac{1}{(1-\gamma)} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)] \end{aligned}$$



Equivalent Form 3 Expressions: Proof

- We have $\nabla V_{\pi_\theta} = (1/(1 - \gamma)) \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we estimate $\widehat{\nabla V_{\pi_\theta}}$?

Equivalent Form 3 Expressions: Proof

- We have $\nabla V_{\pi_\theta} = (1/(1 - \gamma)) \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we estimate $\widehat{\nabla V_{\pi_\theta}}$?
- We will use an equivalent form

Equivalent Form 3 Expressions: Proof

- We have $\nabla V_{\pi_\theta} = (1/(1 - \gamma)) \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)]$
- How do we estimate $\widehat{\nabla V_{\pi_\theta}}$?
- We will use an equivalent form

Proposition

$$\begin{aligned} \nabla V_{\pi_\theta} &= (1/(1 - \gamma)) \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [Q_{\pi_\theta}(s, a) \nabla \log \pi_\theta(s, a)] \\ &= \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_\theta}(s_t, a_t) \nabla \log \pi_\theta(s_t, a_t) \right] \end{aligned}$$

Equivalent Form 3 Expressions: Proof

Proof:



Equivalent Form 3 Expressions: Proof

Proof:

$$\mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right]$$



Equivalent Form 3 Expressions: Proof

Proof:

$$\begin{aligned} & \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \sum_x \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \end{aligned}$$



Equivalent Form 3 Expressions: Proof

Proof:

$$\begin{aligned} & \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \sum_x \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \\ &= \sum_s \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \end{aligned}$$

□

Equivalent Form 3 Expressions: Proof

Proof:

$$\begin{aligned} & \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \sum_x \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \\ &= \sum_s \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \\ &= (1/(1 - \gamma)) \sum_x \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \rho_{\pi_{\theta}}(s, a) \end{aligned}$$

□

Equivalent Form 3 Expressions: Proof

Proof:

$$\begin{aligned} & \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \sum_x \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \\ &= \sum_s \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \\ &= (1/(1 - \gamma)) \sum_x \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \rho_{\pi_{\theta}}(s, a) \\ &= (1/(1 - \gamma)) \mathbb{E}_{(s, a) \sim \rho_{\pi_{\theta}}(\cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)] \end{aligned}$$

□

Equivalent Form 3 Expressions: Proof

Proof:

$$\begin{aligned} & \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_{\theta}}(s_t, a_t) \nabla \log \pi_{\theta}(s_t, a_t) \right] \\ &= \sum_x \sum_a \sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \\ &= \sum_s \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \left(\sum_{t=0}^{\infty} \mathbb{P}(s_t = s, a_t = a) \gamma^t \right) \\ &= (1/(1 - \gamma)) \sum_x \sum_a Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a) \rho_{\pi_{\theta}}(s, a) \\ &= (1/(1 - \gamma)) \mathbb{E}_{(s, a) \sim \rho_{\pi_{\theta}}(\cdot)} [Q_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)] \\ &= \nabla V_{\pi_{\theta}} \end{aligned}$$

□

Variance Reduction using Baseline

- Advantage function of a policy π is defined as

$$A_{\pi}(s, a) = Q_{\pi}(s, a) - V_{\pi}(s)$$

- Note that $\mathbb{E}_{a \sim \pi(s, \cdot)}[A_{\pi}(s, a)] = 0$

Variance Reduction using Baseline

- Advantage function of a policy π is defined as

$$A_{\pi}(s, a) = Q_{\pi}(s, a) - V_{\pi}(s)$$

- Note that $\mathbb{E}_{a \sim \pi(s, \cdot)}[A_{\pi}(s, a)] = 0$

Proposition

$$\nabla V_{\pi_{\theta}} = \frac{1}{(1 - \gamma)} \mathbb{E}_{(s, a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [A_{\pi_{\theta}}(s, a) \nabla \log \pi_{\theta}(s, a)]$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

$$\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)]$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_{\theta}}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [b_{\pi_{\theta}}(s) \nabla \log \pi_{\theta}(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [b_{\pi_{\theta}}(s) \nabla \log \pi_{\theta}(s, a)] \\ = \mathbb{E}_{s \sim \rho_{\pi_{\theta}}(\cdot)} [b_{\pi_{\theta}}(s) \mathbb{E}_{a \sim \pi_{\theta}(s, \cdot)} [\nabla \log \pi_{\theta}(s, a)]] \end{aligned}$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_{\theta}}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [b_{\pi_{\theta}}(s) \nabla \log \pi_{\theta}(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_{\theta}}(\cdot, \cdot)} [b_{\pi_{\theta}}(s) \nabla \log \pi_{\theta}(s, a)] \\ &= \mathbb{E}_{s \sim \rho_{\pi_{\theta}}(\cdot)} [b_{\pi_{\theta}}(s) \mathbb{E}_{a \sim \pi_{\theta}(s, \cdot)} [\nabla \log \pi_{\theta}(s, a)]] \\ &= \mathbb{E}_{s \sim \rho_{\pi_{\theta}}(\cdot)} [b_{\pi_{\theta}}(s) \sum_a \pi_{\theta}(s, a) \nabla \log \pi_{\theta}(s, a)] \end{aligned}$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \mathbb{E}_{a \sim \pi_\theta(s, \cdot)} [\nabla \log \pi_\theta(s, a)]] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \nabla \log \pi_\theta(s, a)] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \nabla \pi_\theta(s, a)] = \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \nabla \sum_a \pi_\theta(s, a)] \end{aligned}$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \mathbb{E}_{a \sim \pi_\theta(s, \cdot)} [\nabla \log \pi_\theta(s, a)]] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \nabla \log \pi_\theta(s, a)] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \nabla \pi_\theta(s, a)] = \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \nabla \sum_a \pi_\theta(s, a)] \\ &= \mathbb{E}_{x \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla(1)] = 0 \end{aligned}$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \mathbb{E}_{a \sim \pi_\theta(s, \cdot)} [\nabla \log \pi_\theta(s, a)]] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \nabla \log \pi_\theta(s, a)] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \nabla \pi_\theta(s, a)] = \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \nabla \sum_a \pi_\theta(s, a)] \\ &= \mathbb{E}_{x \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla(1)] = 0 \end{aligned}$$

So, for any such baseline $b_\pi(s)$, we get

$$\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot)} [(Q_{\pi_\theta}(s, a) - b_{\pi_\theta}(s)) \nabla \log \pi_\theta(s, a)] = \nabla V_{\pi_\theta}$$

Variance Reduction using Baseline

Proof:

Let $b_{\pi_\theta}(s)$ be an action independent baseline.

We will first show that $\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] = 0$.

$$\begin{aligned} \mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla \log \pi_\theta(s, a)] &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \mathbb{E}_{a \sim \pi_\theta(s, \cdot)} [\nabla \log \pi_\theta(s, a)]] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \nabla \log \pi_\theta(s, a)] \\ &= \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \sum_a \nabla \pi_\theta(s, a)] = \mathbb{E}_{s \sim \rho_{\pi_\theta}(\cdot)} [b_{\pi_\theta}(s) \nabla \sum_a \pi_\theta(s, a)] \\ &= \mathbb{E}_{x \sim \rho_{\pi_\theta}(\cdot, \cdot)} [b_{\pi_\theta}(s) \nabla(1)] = 0 \end{aligned}$$

So, for any such baseline $b_\pi(s)$, we get

$$\mathbb{E}_{(s,a) \sim \rho_{\pi_\theta}(\cdot)} [(Q_{\pi_\theta}(s, a) - b_{\pi_\theta}(s)) \nabla \log \pi_\theta(s, a)] = \nabla V_{\pi_\theta}$$

This is in particular true for the baseline $b_{\pi_\theta}(x) = V_{\pi_\theta}(x)$

Policy Gradient Theorem(s)

Theorem

The following as expressions for ∇V_{π_θ} :

$$\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\cdot)} \left[R(\tau) \sum_{t=0}^{\infty} \nabla \log \pi_\theta(s_t, a_t) \right] \quad (1)$$

$$\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t G_t \nabla \log \pi_\theta(s_t, a_t) \right] \quad (2)$$

$$\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t Q_{\pi_\theta}(s_t, a_t) \nabla \log \pi_\theta(s_t, a_t) \right] \quad (3)$$

$$\nabla V_{\pi_\theta} = \mathbb{E}_{\tau \sim P_{\text{traj}, \theta}(\tau)} \left[\sum_{t=0}^{\infty} \gamma^t A_{\pi_\theta}(s_t, a_t) \nabla \log \pi_\theta(s_t, a_t) \right] \quad (4)$$